

Topic 1: Course Introduction, Math Review, and Software

Tian Xie[†]

[†]Singapore Management University and SHUFE

Course Overview

INSTRUCTOR	XIE Tian
EMAIL	tianxie@smu.edu.sg xietian001@hotmail.com
LOCATION	Online via ZOOM (ZOOM info will be posted regularly)
CLASS HOUR	Monday 7:00pm to 10:00pm (ten minutes break between each hour)

Course Contents

- ▶ Unfortunately, ECON.685 Elearn Section is down. I cannot post my course contents on time.

Course Contents

- ▶ Unfortunately, ECON.685 Elearn Section is down. I cannot post my course contents on time.
- ▶ We gonna use **GitHub** instead!!

Course Contents

- ▶ Unfortunately, ECON.685 Elearn Section is down. I cannot post my course contents on time.
- ▶ We gonna use **GitHub** instead!!
- ▶ The link for this course is:

`github.com/xietian001/SMU.ML.Course`

- ▶ **All the course related contents** (outline, schedule, homework, codes, data, etc.) are available via the above link.
- ▶ You don't need to register. Just download the files.
- ▶ Contents are updated on a weekly basis.
- ▶ GitHub is a vastly popular website for codes sharing and project collaborating. Checkout `github.com` for further details.

Purposes of Our Course


- ▶ Know the basics of the machine learning (ML) theory and practice of ML algorithms.
- ▶ Carry out simple empirical exercises using classic ML methods.
- ▶ Summarize and interpret ML results.
- ▶ Discuss the differences between alternative methods commonly used in ML projects.

Assessment Method

- ▶ Assignments (40%)
 - ▶ There will be two assignments handing out.
 - ▶ You are allowed to work in a group of **no more than 5 (including 5)** students and submit one copy of your assignments.
 - ▶ Of course, you can work the assignment just **by yourself**.
 - ▶ You need send the electronic version of your assignments to `tianxie@smu.edu.sg`.
 - ▶ You **must** state all the group members' names clearly on the cover page.
 - ▶ You **must** include the program codes in the assignment.
 - ▶ You can switch groups between the assignments.
- ▶ Class Performance (10%)
- ▶ Final exam (50%)


Recommended Textbook

- ▶ **Statistical Learning from a Regression Perspective (2nd Edition)** by Richard A. Berk.
- ▶ ISBN-13: 978-3319440477
- ▶ ISBN-10: 3319440470




Supplementary Textbook

- ▶ **The Elements of Statistical Learning (2nd Edition)** by Trevor Hastie, Robert Tibshirani, and Jerome Friedman
- ▶ ISBN-13: 978-0387848570
- ▶ ISBN-10: 0387848576



Supplementary Lecture Notes

- ▶ I also uploaded my own lecture notes for you guys.
- ▶ **Machine Learning Lecture Notes** by Tian Xie.
- ▶ For those who don't want have a copy of the textbook, you can read my lecture notes instead.
- ▶ The course slides are abstracted from the notes.
- ▶ We will test contents from the **slides** only.
- ▶ Slides <- My Notes <- Textbook



Contents

- ▶ Splines and Smoothing
- ▶ Classification and Regression Trees
- ▶ Bootstrap and Bagging Tree
- ▶ Random Forest
- ▶ Boosting Tree
- ▶ Support Vector Machine

Machine Learning Concept

A Conventional Introduction

- ▶ Machine learning (ML) is the scientific study of **algorithms and statistical models** that computer systems use to perform a specific task without using explicit instructions, relying on patterns and inference instead.
- ▶ It is seen as a subset of artificial intelligence.
- ▶ The learning process can be categorized as **supervised learning** and **unsupervised learning**.
 - ▶ What is the difference?

Supervised and or Unsupervised?

- ▶ In a typical econometric analysis, we have a pair of \mathbf{X} and \mathbf{y} . For example,

$$\mathbf{y} = \mathbf{X}\beta + \epsilon$$

where

- ▶ \mathbf{X} can be called the regressors, input variables, independent variables, or **features**.
- ▶ \mathbf{y} can be called the regressand, output variable, dependent variable, or **response**.
- ▶ β is the coefficient vector, and ϵ is the error term.
- ▶ **Supervised learning** means, you have both features and the response.
 - ▶ You have input and output. You have a goal to help you decide/evaluate.
- ▶ **Unsupervised learning** means, you only have features.
 - ▶ You only have \mathbf{X} .
 - ▶ You try to **learn the pattern** lurking inside of a data set.
- ▶ **Most of the economic problems we study require supervised learning.**

Supervised Learning

- ▶ For supervised learning, depending on **type** of the response variable, the problem we study is either a **classification** problem or a **regression** problem.

Supervised Learning

- ▶ For supervised learning, depending on **type** of the response variable, the problem we study is either a **classification** problem or a **regression** problem.
 - ▶ analyze pets: dog, cat, bird, fish...
 - ▶ forecast stock price using capitalization, liquidity, age of CEO...

Supervised Learning

- ▶ For supervised learning, depending on **type** of the response variable, the problem we study is either a **classification** problem or a **regression** problem.
 - ▶ analyze pets: dog, cat, bird, fish...
 - ▶ forecast stock price using capitalization, liquidity, age of CEO...
- ▶ Classification requires **categorical** responses and Regression requires **numerical** responses.
 - ▶ Which problem is more frequently encountered in economics?

Supervised Learning

- ▶ For supervised learning, depending on **type** of the response variable, the problem we study is either a **classification** problem or a **regression** problem.
 - ▶ analyze pets: dog, cat, bird, fish...
 - ▶ forecast stock price using capitalization, liquidity, age of CEO...
- ▶ Classification requires **categorical** responses and Regression requires **numerical** responses.
 - ▶ Which problem is more frequently encountered in economics?
 - ▶ Regression analysis is more popular in economics and finance.

Nonlinearity and Flexibility

- ▶ Huge hype about machine learning in **Economics and Finance** now.
- ▶ Many people apply **fancy** ML algorithms to economic problem **brutally** without even knowing the reason and logic.
- ▶ Remember that we are studying **Economics and Finance**. There has to be some **motivation**.
- ▶ Of course, every data is unique. However, Economics and Finance data do have universal **patterns**.
 - ▶ For example, stocks prices are very hard to forecast, but stock volatilities are easy to predict.
 - ▶ It is common that certain algorithms have better performance than others.

Nonlinearity and Flexibility

- ▶ Huge hype about machine learning in **Economics and Finance** now.
- ▶ Many people apply **fancy** ML algorithms to economic problem **brutally** without even knowing the reason and logic.
- ▶ Remember that we are studying **Economics and Finance**. There has to be some **motivation**.
- ▶ Of course, every data is unique. However, Economics and Finance data do have universal **patterns**.
 - ▶ For example, stocks prices are very hard to forecast, but stock volatilities are easy to predict.
 - ▶ It is common that certain algorithms have better performance than others.
- ▶ Many ML algorithms are **nonlinear** and **flexible**. They break the barriers of **linearity** and **parametric** formulation.
 - ▶ That is why they have good performance.
 - ▶ But say the data is super linear, a nonlinear algorithm shouldn't have a huge advantage.

Coding

The Role of Coding

- ▶ Coding is very **important** in studying ML.
 - ▶ Coding helps you better understand the contents and forces you to pay attention to shuttle details.

The Role of Coding

- ▶ Coding is very **important** in studying ML.
 - ▶ Coding helps you better understand the contents and forces you to pay attention to shuttle details.
- ▶ But our course is not called “ML in R or ML in MATLAB”.
 - ▶ The primary concern of the course is not coding.
 - ▶ We will **NOT** test your coding skills in the final.
- ▶ Learning ML without coding is like learning swimming without getting wet.
- ▶ Following the Dean’s “suggestion”, we mainly use **R** in this course to demonstrate coding and estimation, therefore, you are recommended to follow our choice of software.
 - ▶ You are free to use whatever software you like, for example, Eviews, Stata, Matlab, R, Python, Java, C, C++, or even MS Excel, as long as you can deliver qualified course work.

R and RStudio - The Old-school Way

- ▶ To use R, you need to the **R source files first**.
 - ▶ You can obtain the files from <https://www.r-project.org/>.
 - ▶ It has many different versions that can generate various instability/incompatibility problems.
 - ▶ Have fun!
- ▶ Then, you need a good **R composer with nice UI**. The most popular one is **RStudio**.
 - ▶ You can obtain the free open source version from <https://rstudio.com/products/rstudio/download/>
- ▶ They are free and small size (less than 300M in total).
- ▶ You can install them in your own computer or in a flashdrive.
 - ▶ For flashdrive installation, you can plug-in and use immediately.
- ▶ **But we are not gonna do any of the above in this course.**

RStudio Cloud

- ▶ In this course, you are highly recommend to use the **RStudio Cloud** to learn R syntax, practice exercises, and do homeworks.
 - ▶ Visit the link:
rstudio.cloud
 - ▶ **Register a free account** and start coding!
- ▶ Cloud computing has lots of merits:
 - ▶ No installation needed! Simply open a browser and stay online!
 - ▶ No instability or incompatibility.
 - ▶ The cloud records every steps of your coding process, so you never lose your codes, data, etc.
- ▶ Perhaps, the only drawback of cloud computing is that you have to stay online.

Console Window

- ▶ In the **Console window**, R responds to any input immediately, like a calculator. Try

```
> 1+1  
[1] 2
```

Notice that R immediately responds to your input and [1] implies the results are listed in the first row.

- ▶ R didn't record the result, since you didn't assign a variable. You can assign a variable to complicate calculation for re-use purpose. Try

```
> x = 1+1
```

Checkout the **Environment window**. You may notice a variable `x` with value 2.

- ▶ ^(?)Try `y <- 1+1`. What is the difference between `=` and `<-`?

Functions

- ▶ R can do way more than calculators. Try `rnorm(3)`.

```
> rnorm(3)
[1]  2.2461109  0.6867319 -0.7039494
```

- ▶ I am sure your results are different from mine. ^(?)Do you know why?
- ▶ To fully understand this command, its meaning, function, syntax, etc. Use `help(rnorm)` or simply type `?rnorm`.
- ▶ Its information is presented in the **help window**. Try to digest its meaning.
- ▶ ^(?)Generate 20 random results from $N(10, 25)$.
- ▶ If your Console is messy, try `Ctrl+L` to clean the window.


Exercise

1. Let us generate 10 random variables from standard normal distribution and compute their **mean** and **variance**.
**Note that you may need the functions: `mean`, `var`.*
2. Let us generate 100 random variables from standard normal distribution and compute their **mean** and **variance**.
3. Let us generate 10000 random variables from standard normal distribution and compute their **mean** and **variance**.

Notice any pattern?

Script Files

- ▶ Using the Console window to execute commands is rather inefficient.
- ▶ Like many other programmable software, we can use create a **script file** that consists of multiple lines of command and execute them **in sequence**.
- ▶ Click this icon to create an empty script and save this file in a designated **location** with proper **file name**.



- ▶ We can write the following lines to the script.

```
# Mean and Variance of normal RVs
x1 = rnorm(10)
x2 = rnorm(100)
x3 = rnorm(10000)
m1 = mean(x1)
m2 = mean(x2)
m3 = mean(x3)
v1 = var(x1)
v2 = var(x2)
v3 = var(x3)
```

- ▶ Select the lines you want to execute and click Run or use Ctrl+Enter. You should notice the new results in the **Environment Window**.
- ▶ You can use # to add **comments**. Contents after # are not executed.

Loops

- ▶ Now let us consider the following exercise: generate 100, 200, 300, ..., 100000 random variables from $N(1, 1.5)$ and compute their means and variances.
- ▶ If you manually type up 100, 200, 300, ..., 100000, it will take forever to complete.
- ▶ Command `for` can repeat a pre-defined process multiple times. We usually refer this procedure as **loops**. The syntax of `for` is

```
for (indicator in sequence){  
  code  
}
```

where **indicator** can be any parameters, `i`, `j`, ...
sequence represents a sequence of data
code can any estimating function you design

- ▶ Here is a demo code:

```
# use loop to obtain mean and variance
MEAN = 0;
VAR = 1;
n = seq(100,100000,by=100);
for (i in 1:length(n)){
  x = rnorm(n[i],mean=1,sd=sqrt(1.5))
  MEAN[i] = mean(x)
  VAR[i] = var(x)
}
```


Figures

- It is more intuitive to plot variables MEAN and VAR in **figures**.

```
# plot MEAN and VAR separately
plot(n,MEAN,col="blue",type="l",lwd=1,
      xlab="Sample Size", ylab="Value")
plot(n,VAR,col="red",type="l",lty="dashed",lwd=1,
      xlab="Sample Size", ylab="Value")
```

- Or you can plot both lines in the **same figure**.

```
# plot MEAN and VAR together
plot(n,MEAN,col="blue",type="l",lwd=1,
      xlab="Sample Size", ylab="Value",ylim=c(0.9,2))
lines(n,VAR,col="red",lty="dashed",lwd=1,
      xlab="Sample Size", ylab="Value")
legend(100,2,legend=c("Mean","Variance"),
      col=c("Blue","Red"),lty=1:2)
```

Exercise

- ▶ Plot the PDF of $N(0, 1)$.
- ▶ Plot the PDF of t_2^2 .
- ▶ Plot the CDF of t_{25} .
- ▶ Merge all three plots in one figure.
- ▶ You may need the command `dnorm`, `dt`, `pt`.

Answers

```
# plot the distribution
x = seq(-4,4,by=0.1)
y1 = dnorm(x,mean=0,sd=1)
y2 = dt(x,df=2)
y3 = pt(x,df=25)
plot(x,y1,type='l',lwd=1,col="blue",
      xlab="x",ylab="y",ylim=c(0,1))
lines(x,y2,lty="dashed",lwd=1,col="red")
lines(x,y3,lty="dotted",lwd=1,col="black")
legend(-4,1,legend=c("PDF of N(0,1)","PDF of t(2)","CDF of t(25)"),
      col=c("Blue","Red","Black"),lty=1:3)
```

Import Data

- ▶ In practice, it is quite common to performance analysis on **given data set**.
- ▶ Here we use the `movie.csv` data file to demonstrate.
 - ▶ First, you need to download the data `movie.csv` from github.com/xietian001/SMU.ML.Course.
 - ▶ Then, you need to **upload** the data to the **cloud**.
 - ▶ This data consists of 94 movies with their **open box office** and related variables. (?)What do you think determine a movie's sales?
- ▶ We can use `read.csv` command to import the data. We need tell R the **exact location** of the file.

```
# import movie data
LOC = "/cloud/project/movie.csv"
dat = read.csv(LOC,header=TRUE)
summary(dat)
```

- ▶ (?)What is the functionality of `summary()`?
- ▶ The variable `dat` is a stored in a **list** format.
 - ▶ Click the `dat` to see its contents.
 - ▶ To access each element, for example, `OpenBox`, you can use `dat$OpenBox`.

Exercise

- ▶ Plot **open box office** against **budgets** and add a 45° line. What can you conclude?


Exercise

- Plot **open box office** against **budgets** and add a 45° line. What can you conclude?

```
# plot open box office against budget
x = dat$Budget
y = dat$OpenBox
plot(x,y,col="blue",pch = 16,xlim=c(0,16),ylim=c(0,16),
      xlab="Budget",ylab="Open Box Office")
lines(c(0,16),c(0,16),type="l",col="black")
```

Movie Plot


- Here is the plot using movie data.



How to Code?

- ▶ There is **no simple** answer.
- ▶ Remember, you are not professional programmer.
- ▶ You (probably) will not code for a living.

Figure: My Marking Standard



Math Review

Linear algebra

- ▶ We denote the set of all n -tuples of real numbers by \mathbb{R}^n . The set of n -tuples of nonnegative real numbers is denoted by \mathbb{R}_+^n .
- ▶ The elements of these sets will be referred to as points or **vectors**.
- ▶ If $x = (x_1, \dots, x_n)$ is a vector, we denote then its i^{th} component is x_i .
- ▶ We can add two vectors by adding their components:
$$x + y = (x_1 + y_1, \dots, x_n + y_n).$$
- ▶ We can perform scalar multiplication on a vector by multiplying every component by a fixed real number t : $tx = (tx_1, \dots, tx_n)$.

- ▶ A vector x is a linear combination of a set of n vectors A if $x = \sum_{i=1}^n t_i y_i$, where $y_i \in A$ and the t_i 's are scalars.
- ▶ A set A of n vectors is linearly independent if there is no set (t_i, x_i) , with some $t_i \neq 0$ and $x_i \in A$, such that $\sum_{i=1}^n t_i x_i = 0$.
- ▶ An equivalent definition is that no vector in A can be represented as a linear combination of vectors in A .
- ▶ Given two vectors their **inner product** is given by $xy = \sum_i x_i y_i$. The norm of a vector x is denoted by $|x|$ and defined by $|x| = \sqrt{xx}$.
- ▶ Note that by the Pythagorean theorem, the norm of x is the distance of the point x from the origin; that is, it is the length of the vector x .
- ▶ If $xy = 0$, then x and y are said to be orthogonal.

- ▶ Let θ be the angle between x and y . It is clear $t|x| = |y|\cos\theta$. Moreover, $xy = |x||y|\cos\theta$.
- ▶ We can consider maps from \mathbb{R}^n to \mathbb{R}^m that send vectors into vectors. We denote such maps by $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$.
- ▶ A map is a **linear** function if $f(tx + sy) = tf(x) + sf(y)$ for all scalars s and t and vectors x and y .
- ▶ If f is a linear function to \mathbb{R}^1 , we call it a linear functional. If p is a linear functional we can represent it by a vector $p = (p_1, \dots, p_n)$, and write $p(x) = px$.
- ▶ A set of points of form $H(p, a) = \{x : px = a\}$ is called a **hyperplane**.

Definite and semidefinite matrices

- ▶ Let A be a symmetric square matrix. Then if we post-multiply A by some vector x and pre-multiply it by the **transpose of the** same vector x , we have a quadratic form.

$$(x_1 \quad x_2) \begin{pmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = a_{11}x_1^2 + (a_{21} + a_{12})x_1x_2 + a_{22}x_2^2.$$

- ▶ Suppose that A is the identity matrix. In this case it is not hard to see that whatever the values of x_1 and x_2 , the quadratic form must be nonnegative.
- ▶ In fact, if x_1 and x_2 are not both zero, xAx^T will be strictly positive. The identity matrix is an example of a positive definite matrix.

- ▶ **Definite matrices.** A square matrix A is:
 - (a) positive definite if $x^T A x > 0$ for all $x \neq 0$;
 - (b) negative definite if $x^T A x < 0$ for all $x \neq 0$;
 - (c) positive semidefinite if $x^T A x \geq 0$ for all x ;
 - (d) negative semidefinite if $x^T A x \leq 0$ for all x .
- ▶ We say A is positive definite subject to constraint $b x = 0$ if $x^T A x > 0$ for all $x \neq 0$ such that $b x = 0$. The other definitions extend to the constrained case in a natural manner.

- ▶ If a matrix is positive semidefinite, then it must have nonnegative diagonal terms.
- ▶ The minor matrices of a matrix A are the matrices formed by eliminating k columns and the same numbered k rows. The naturally ordered or nested principal minor matrices of A are the minor matrices given by

$$a_{11} \quad \begin{pmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{pmatrix} \quad \begin{pmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{pmatrix}$$

- ▶ The minor determinants or minors of a matrix are the determinants of the minors. We denote the determinant of a matrix A by $\det A$ or $|A|$.

Cramer's rule

- ▶ Here is a convenient rule for solving linear systems of equations of the form

$$\begin{pmatrix} a_{11} & \cdots & a_{1n} \\ \vdots & & \vdots \\ a_{n1} & \cdots & a_{nn} \end{pmatrix} \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix} = \begin{pmatrix} b_1 \\ \vdots \\ b_n \end{pmatrix}$$

We can write this system more conveniently as $Ax = b$.

- ▶ **Cramer's rule.** To find the component x_i of the solution vector to this system of linear equations, replace the i^{th} column of the matrix A with the column vector b to form a matrix A_i . Then x_i is the determinant of A_i , divided by the determinant of A :

$$x_i = \frac{|A_i|}{|A|}.$$

Calculus

- ▶ Given a function $f : \mathbb{R} \rightarrow \mathbb{R}$, we define its derivative at a point x^* by

$$\frac{df(x^*)}{dx} = \lim_{t \rightarrow 0} \frac{f(x^* + t) - f(x^*)}{t}$$

if that limit exists.

- ▶ The derivative $df(x^*)/dx$ is also denoted by $f'(x^*)$. If the derivative of f exists at x^* , we say that f is **differentiable** at x^* .
- ▶ Consider the linear function $F(t)$ defined by

$$F(t) = f(x^*) + f'(x^*)t.$$

This is a good approximation to f near x^* since

$$\lim_{t \rightarrow 0} \frac{f(x^* + t) - F(t)}{t} = \lim_{t \rightarrow 0} \frac{f(x^* + t) - f(x^*) - f'(x^*)t}{t} = 0.$$

- In the same way, given an arbitrary function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, we can define its derivative at x^* , $D_f(x^*)$, as being that linear map from \mathbb{R}^n to \mathbb{R}^m that approximates f close to x^* in the sense that

$$\lim_{|t| \rightarrow 0} \frac{|f(x^* + t) - f(x^*) - D_f(x^*)t|}{|t|} = 0.$$

- We use norm signs since both the numerator and denominator are vectors. The map $f(x^*) + D_f(x^*)t$ is a good approximation to f at x^* in the sense that for small vectors t ,

$$f(x^* + t) \approx f(x^*) + D_f(x^*)t.$$

- ▶ Given a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, we can also define the **partial derivatives** of f with respect to x_i evaluated at x^* .
- ▶ To do this, we hold all components fixed except for the i^{th} component, so that f is only a function of x_i , and calculate the ordinary one-dimensional derivative.
- ▶ We denote the partial derivative of f with respect to x_i evaluated at x^* by $\partial f(x^*)/\partial x_i$.
- ▶ Since $D_f(x^*)$ is a linear transformation, we can represent it by a matrix, which turns out to be

$$D_f(x^*) = \begin{pmatrix} \frac{\partial f_1(x^*)}{\partial x_1} & \dots & \frac{\partial f_1(x^*)}{\partial x_n} \\ \vdots & & \vdots \\ \frac{\partial f_m(x^*)}{\partial x_1} & \dots & \frac{\partial f_m(x^*)}{\partial x_n} \end{pmatrix}.$$

The matrix representing $D_f(x)$ is called the **Jacobian matrix** of f at x^* . We will often work with functions from \mathbb{R}^n to \mathbb{R} in which case $D_f(x^*)$ will be an $n - by - 1$ matrix, which is simply a vector.

Higher-order derivatives

- ▶ If we have a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, the **Hessian matrix** of that function is the matrix of mixed partial derivatives

$$D^2 f(x) = \frac{\partial^2 f(x)}{\partial x_i \partial x_j}.$$

Note that $D^2 f(x)$ is a symmetric matrix.

- ▶ Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a differentiable function and let x and y be two vectors in \mathbb{R}^n . Then it can be shown that

$$\begin{aligned} f(y) &= f(x) + D_f(z)(y - x) \\ f(y) &= f(x) + D_f(x)(y - x) + \frac{1}{2}(y - x)^\top D^2 f(w)(y - x). \end{aligned}$$

where z and w are points on the line segment between x and y . These expressions are called **Taylor series expansions** of f at x .

- If x and y are close together and the derivative functions are continuous, then $Df(z)$ and $D^2f(w)$ are approximately equal to $Df(x)$ and $D^2f(x)$, respectively. We therefore often write the Taylor series expansions as

$$f(y) = f(x) + D_f(x)(y - x)$$

$$f(y) = f(x) + D_f(x)(y - x) + \frac{1}{2}(y - x)^\top D^2f(x)(y - x).$$

Analysis

- ▶ Given a vector x in \mathbb{R}^n and a positive real number e , we define an **open ball** of radius e at x as $B_e(x) = \{y \in \mathbb{R}^n : |y - x| < e\}$.
- ▶ A set of points A is a **open set** if for every x in A there is some $B_e(x)$ which is contained in A .
- ▶ If x is in an arbitrary set and there exists an $e > 0$ such that $B_e(x)$ is in A , then x is said to be in the **interior** of A .
- ▶ The complement of a set A in \mathbb{R}^n consists of all the points in \mathbb{R}^n that are not in A ; it is denoted by $\mathbb{R}^n \setminus A$.
- ▶ A set is a **closed set** if $\mathbb{R}^n \setminus A$ is an open set. A set A is bounded if there is some x in A and some $e > 0$ such that A is contained in $B_e(x)$. If a nonempty set in \mathbb{R}^n is both closed and bounded, it is called **compact**.
- ▶ A infinite **sequence** in \mathbb{R}^n , $(x^i) = (x^1, x^2, \dots)$ is just an infinite set of points, one point for each positive integer.

- ▶ A sequence (x^i) is said to converge to a point x^* if for every $\epsilon > 0$, there is an integer m such that, for all $i > m$, x^i is in $B_\epsilon(x^*)$. We sometimes say that x^i gets arbitrarily close to x^* . We also say that x^* is the **limit** of the sequence (x^i) and write $\lim_{i \rightarrow \infty} x^i = x^*$. If a sequence converges to a point, we call it a **convergent sequence**.
- ▶ **Closed set**. A is a closed set if every convergent sequence in A converges to a point in A .
- ▶ **Compact set**. If A is a compact set, then every sequence in A has a convergent subsequence.
- ▶ A function $f(x)$ is continuous at x^* if for every sequence (x^i) that converges to x^* , we have the sequence $(f(x^i))$ converging to $f(x^*)$. A function that is continuous at every point in its domain is called a **continuous function**.

Random Variables and Probability Distributions

- ▶ **Random variables.** A random variable is a numerical summary of a random outcome.
 - ▶ **Discrete** random variable: takes on only a discrete set of values (toss coins, roll dice,...). **Countable output.**
 - ▶ **Continuous** random variable takes on a continuum of possible values (stock return, GDP growth rate, ...). **Non-countable output.**
- ▶ The **probability** of an outcome is the proportion of the time that the outcome is **expected** to occur in the long run.
 - ▶ ^(?)Question: what is difference between probability and frequency?

Random Variables and Probability Distributions

- ▶ **Random variables.** A random variable is a numerical summary of a random outcome.
 - ▶ **Discrete** random variable: takes on only a discrete set of values (toss coins, roll dice,...). **Countable output.**
 - ▶ **Continuous** random variable takes on a continuum of possible values (stock return, GDP growth rate, ...). **Non-countable output.**
- ▶ The **probability** of an outcome is the proportion of the time that the outcome is **expected** to occur in the long run.
 - ▶ (?) Question: what is difference between probability and frequency?
 - ▶ (?) or say, what is difference between 0% and never gonna happen?

Random Variables and Probability Distributions

- ▶ **Random variables.** A random variable is a numerical summary of a random outcome.
 - ▶ **Discrete** random variable: takes on only a discrete set of values (toss coins, roll dice,...). **Countable output.**
 - ▶ **Continuous** random variable takes on a continuum of possible values (stock return, GDP growth rate, ...). **Non-countable output.**
- ▶ The **probability** of an outcome is the proportion of the time that the outcome is **expected** to occur in the long run.
 - ▶ ^(?)Question: what is difference between probability and frequency?
^(?)or say, what is difference between 0% and never gonna happen?
In probability space, they are 0 and \emptyset

- ▶ **Probability distribution**. The probability distribution of a discrete random variable is the list of all possible values of the variable and the probability that each value will occur. These probabilities sum to 1.
- ▶ The **cumulative probability distribution** is the probability that the random variable is less than or equal to a particular value.
 - ▶ It is also referred to as a **cumulative distribution function (CDF)**.

Table: Probability of Donald Trump Tweeting X Times Per Day

	Number of Tweets					
	0	1	2	3	4	5
Probability Distribution	0.1	0.2	0.2	0.3	0.1	0.1
Cumulative Probability Distribution	0.1	0.3	0.5	0.8	0.9	1.0


Bernoulli distribution

- ▶ A **binary** random variable is called a **Bernoulli** random variable, and its probability distribution is called the Bernoulli distribution.
- ▶ The outcome of a Bernoulli random variable can only take two values.
 - ▶ For example, 1 and 0, True and False, etc.
- ▶ For a Bernoulli random variable X with two outcomes 1 and 0, the probability distribution can be expressed as

$$f(X) = \begin{cases} 1 & \text{with probability } p \\ 0 & \text{with probability } (1 - p) \end{cases}$$

Continuous Random Variable

- ▶ A **continuous** random variable can have **infinite** number of outcomes.
- ▶ It can be constrained within a certain range or **unconstrained**.
- ▶ A typical example is the standard **Gaussian** random variable with probability **density** function (PDF) and CDF shown below:



Expectation

- ▶ The **expectation** of a random variable X measures its long-run average value. It is the sum of the any outcome times its respective probability.

$$\mathbb{E}(X) = \sum_{i=1}^N X_i \cdot p_i,$$

where X_i is one outcome and p_i is the associated probability.

- ▶ Recall the Tweeting example:

	Number of Tweets					
	0	1	2	3	4	5
Probability	0.1	0.2	0.2	0.3	0.1	0.1

- (?) What is the expectation of Trump's tweets per day?
- ▶ If a is a constant, $\mathbb{E}(aX) = a\mathbb{E}(X)$.
- ▶ We also have $\mathbb{E}(X + Y) = \mathbb{E}(X) + \mathbb{E}(Y)$ for any X and Y .

Variance

- ▶ The **variance** measures the **dispersion** or the **spread** of a probability distribution:

$$\text{Var}(X) = \mathbb{E} \left[(X - \mathbb{E}(X))^2 \right] = \sum_{i=1}^N (X_i - \mathbb{E}(X))^2 \cdot p_i.$$

- ▶ ^(?)What is the variance of Trump's daily tweets?
- ▶ **Standard deviation** is simply the square-root of variance.
- ▶ If a and b are constant, $\text{Var}(aX) = a^2 \text{Var}(X)$ and $\text{Var}(X + b) = \text{Var}(X)$.
- ▶ ^(?)Skewness and Kurtosis.

Joint Distribution

- ▶ The **joint** probability distribution of **multiple** random variables is the probability that the random variables **simultaneously** take on certain values.
- ▶ Let us now think about the two **sentimental** stage of D. Trump and relate to his tweeting behavior.
- ▶ Assume Trump's sentiment follows a Bernoulli distribution with two outcomes: **Rage** and **Calm**.



Rage Vs. Calm?

Marginal Distribution

- ▶ Now, Trump's tweeting behavior is measured by two random variables

Table: Joint Distribution of Donald Trump's Tweeting and Sentiment

Sentiment	Number of Tweets						Total
	0	1	2	3	4	5	
Rage	0.1	0.15	0.1	0.2	0.05	0.1	0.7
Calm	0	0.05	0.1	0.1	0.05	0	0.3
Total	0.1	0.2	0.2	0.3	0.1	0.1	1.0

- ▶ **Marginal probability distribution** is the sum of individual probabilities associated with one specific outcome of a certain random variable.
 - ▶ For example, $\Pr(\text{Rage}) = 0.1 + 0.15 + \dots + 0.1 = 0.7$.

Conditional Distribution

- ▶ **Conditional distribution** is the distribution of a random variable Y conditional on another random variable X taking on a specific value, usually denoted as $\Pr(Y = y|X = x)$.
- ▶ The conditional probability can be estimated by

$$\Pr(X = x|Y = y) = \frac{\Pr(X = x, Y = y)}{\Pr(Y = y)}.$$

- ▶ **Conditional Expectation** and **Conditional Variance** are

$$\mathbb{E}(X|Y = y) = \sum_{i=1}^N X_i \cdot \Pr(X = X_i|Y = y),$$

$$\text{Var}(X|Y = y) = \sum_{i=1}^N (X_i - \mathbb{E}(X|Y = y))^2 \cdot \Pr(X = X_i|Y = y),$$

Independence, Covariance, and Correlation

- ▶ Two random variables X and Y are **independent**, if knowing the value of one of the variables provides **no information** about the other.
 - ▶ If X and Y are independent, for all values of x and y , we have $\Pr(Y = y|X = x) = \Pr(Y = y)$.
- ▶ One measure of the extent to which two random variables move together is their **covariance**.
 - ▶ Specifically,
$$\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}(X))(Y - \mathbb{E}(Y))] = \mathbb{E}(XY) - \mathbb{E}(X)\mathbb{E}(Y).$$
- ▶ The correlation is an **alternative measure** of dependence between X and Y that follows unity.
 - ▶ Specifically,

$$\text{Corr}(X, Y) = \frac{\text{Cov}(X, Y)}{\sqrt{\text{Var}(X) \cdot \text{Var}(Y)}}$$

Independence, Covariance, and Correlation

- ▶ Independence implies both Covariance and Correlation equal 0.
 - ▶ ^(?) *vice versa*?
 - ▶ Try to the Covariance of $\mathbb{E}(X) = \mathbb{E}(X^3) = 0$, $Y = X^2$.
- ▶ We usually denote the variance as σ^2 , for example $\text{Var}(X) = \sigma_X^2$.
 - ▶ $\text{Cov}(X, Y) = \sigma_{XY}$
 - ▶ Standard Deviation of X is σ_X
- ▶ $\text{Var}(aX + bY) = a^2\sigma_X^2 + b^2\sigma_Y^2 + 2ab \cdot \sigma_{XY}$.
- ▶ By definition, $|\text{corr}(X, Y)| \leq 1$.
- ▶ ^(?) If X and Y are independent, what is $\text{Var}(X + Y)$?